

Project ML Fall 2015. Appendix 2.1

Anastasiya Yarygina Udovenko, Manuel Aragonés Mora, Andrés Ponce de Leon

December 7, 2015

```
rm(list=ls())  
setwd("~/Dropbox/MPP/ML/project")
```

II.1 Building Predictive Model for Poverty Threshold

```
# load required packages, set up parallel computation
```

```
library(recommenderlab)  
library(ggplot2)  
library(caret)  
library(doParallel)  
library(randomForest)  
library(gbm)  
require(ROCR)
```

```
cl <- makeCluster(detectCores() - 2)  
clusterEvalQ(cl, library(foreach))  
registerDoParallel(cl) # register this cluster
```

Set seed

```
set.seed(99)  
  
# read data  
data<- read.csv("dataProject.csv", sep="," , header=TRUE)  
dim(data)
```

```
## [1] 3146 633
```

Some data cleaning

```
# remove last three rows:  
n<-dim(data)[1]  
data<-data[1:(n-3),]  
  
# get mean of Poverty.Percent:  
summary(data$Poverty.Percent)[4]
```

```
## Mean  
## 17.22
```

```

# generate new factor variable poverty
data$poverty <- ifelse(data$Poverty.Percent<= 17.22,
                      c("below"), c("above"))

data$poverty<- as.factor(data$poverty)

# remove regressors that we will not use: est variables
data2<- data[,-grep("est", colnames(data))]

# remove more regressors that we will not use
data3<- data2[, -which(names(data) %in% c("X.2", "X.1", "X", "NAME", "STATE_NAME",
                                         "STATE_FIPS", "CNTY_FIPS", "FIPS", "Poverty.Percent",
                                         "Median.Household.Income"))]
data4<- data3[,-grep("^X",colnames(data3))]

```

Our first outcome variable is “poverty threshold” it takes value 1 if county’s poverty index is below national average and zero otherwise

```

# define y as outcome variable
y<- data4$poverty

```

Some variables are imported in formats that are not suitable for our analysis. We transform them accordingly

```

# convert socio-economic indicators into integers
cols <- data4[, -which(names(data4) %in% c("poverty"))]
cols<- cols[, -grep("rca", colnames(cols))]
cols <- data.frame(apply(cols, 2, as.integer))

# convert index of competitiveness into factors
cols1<- data4[,grep("rca", colnames(data4))]
cols1<- data.frame(apply(cols1, 2, as.factor))

#bind socio-economic indicators and index of competitiveness
#into dataset of analysis:

data_new<- cbind(cols, cols1)

# add the dependent variable
data_new$y <- y

```

Models training and fitting

```

# train and fit predictive models for the outcome "poverty" index:

# create data partition: %60 train and 40% test:
inTrain<- createDataPartition(y=data_new$y,
                              p=0.60, list=FALSE)

trainDf<- data_new[inTrain,]
testDf<- data_new[-inTrain,]

pnm="Counties"

```

We take advantage of the R code provided with the lecture notes (week4), to construct “loss” and “lift” functions that we use to choose the best model specification fitting our dataset.

```
#deviance loss function
lossf = function(y,phat,wht=0.0000001) {
  #y should be 0/1
  #wht shrinks probs in phat towards .5, don't log 0!
  if(is.factor(y)) y = as.numeric(y)-1
  phat = (1-wht)*phat + wht*.5
  py = ifelse(y==1,phat,1-phat)
  return(-2*sum(log(py)))
}

#lift function
liftf = function(yl,phatL,dopl=TRUE) {
  if(is.factor(yl)) yl = as.numeric(yl)-1
  oo = order(-phatL)
  sy = cumsum(yl[oo])/sum(yl==1)
  if(dopl) {
    ii = (1:length(sy))/length(sy)
    plot(ii,sy,type='l',lwd=2,col='blue',xlab='% tried',ylab='% of successes',cex.lab=2)
    abline(0,1,lty=2)
  }
  return(sy)
}

# I initialize the list where we store the results
phatL = list()
```

1. Fit Logit Model

```
# set up
phatL$logit = matrix(0.0,nrow(testDf),1)

# fit logit

lgfit = glm(y~.,trainDf,family=binomial)

# predict
phat = predict(lgfit,testDf,type="response")

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

# store prediction
phatL$logit = matrix(phat,ncol=1)
```

2. Fit Random Forest Models

```

# set up
p=ncol(trainDf)-1
mtryv = c(p,sqrt(p))
ntreev = c(500,1000)
setrf = expand.grid(mtryv,ntreev)
colnames(setrf)=c("mtry","ntree")
phatL$rf = matrix(0.0,nrow(testDf),nrow(setrf))

# train and fit

for(i in 1:nrow(setrf)) {
  cat("on randomForest fit ",i,"\n")
  print(setrf[i,])

  frf = randomForest(y~.,data=trainDf,mtry=setrf[i,1],ntree=setrf[i,2])
  phat = predict(frf,newdata=testDf,type="prob")[,2]

  phatL$rf[,i]=phat # store results in matrix phatL
}

```

3. Fit Boosting Models

```

idv = c(2,4)
ntv = c(1000,5000)
shv = c(.1,.01)
setboost = expand.grid(idv,ntv,shv)
colnames(setboost) = c("tdepth","ntree","shrink")
phatL$boost = matrix(0.0,nrow(testDf),nrow(setboost))

trainDfB = trainDf; trainDfB$y = as.numeric(trainDfB$y)-1
testDfB = testDf; testDfB$y = as.numeric(testDfB$y)-1

# train and fit

for(i in 1:nrow(setboost)) {
  cat("on boosting fit ",i,"\n")
  print(setboost[i,])

  fboost = gbm(y~.,data=trainDfB,distribution="bernoulli",
               n.trees=setboost[i,2],interaction.depth=setboost[i,1],shrinkage=setboost[i,3])
  phat = predict(fboost,newdata=testDfB,n.trees=setboost[i,2],type="response")

  phatL$boost[,i] = phat # store results in the phatL matrix
}

```

4. Compute and plot deviance loss for Logit, RF and Boosting models

```

lossL = list()
nmethod = length(phatL)

```

```

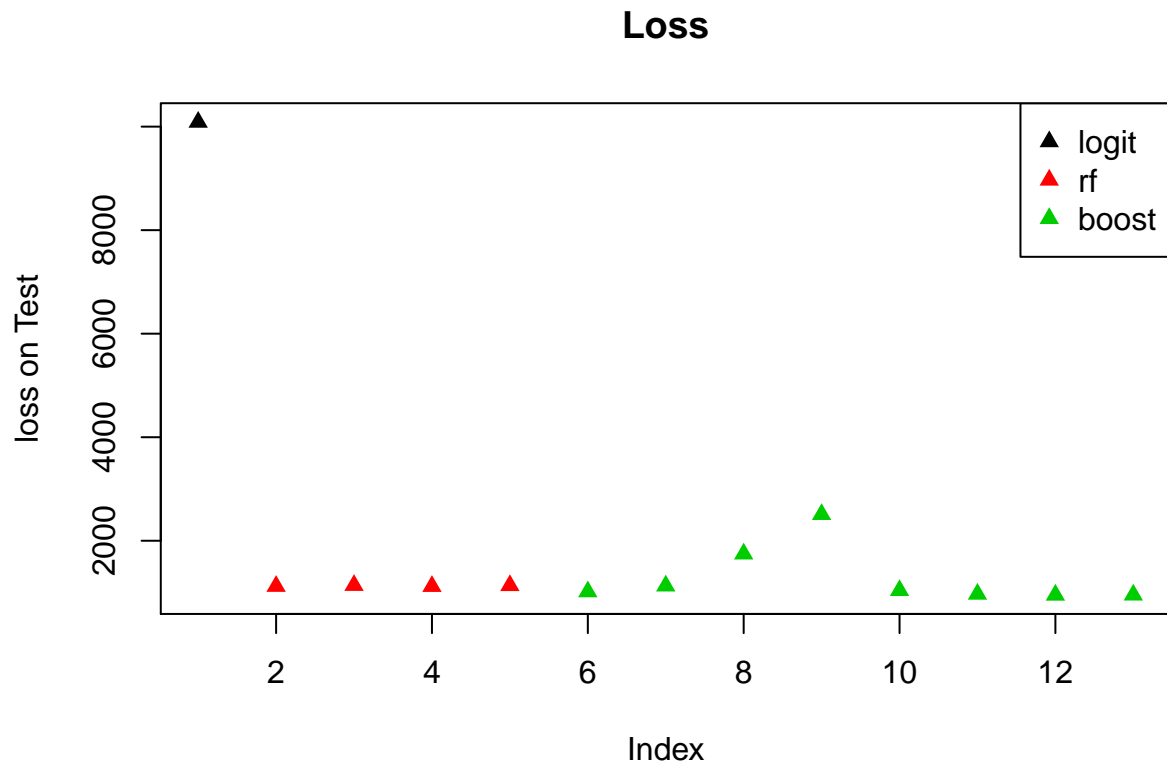
for(i in 1:nmethod) {
  nrun = ncol(phatL[[i]])
  lvec = rep(0,nrun)
  print(nrun)
  for(j in 1:nrun) lvec[j] = lossf(testDf$y,phatL[[i]][,j])
  lossL[[i]]=lvec; names(lossL)[i] = names(phatL)[i]
}

```

```

lossv = unlist(lossL)
par(mfrow=c(1,1))
plot(lossv,ylab="loss on Test",type="n", main = "Loss")
nloss=0
for(i in 1:nmethod) {
  ii = nloss + 1:ncol(phatL[[i]])
  points(ii,lossv[ii],col=i,pch=17)
  nloss = nloss + ncol(phatL[[i]])
}
legend("topright",legend=names(phatL),col=1:nmethod,pch=rep(17,nmethod))

```



From this figure we can see that Logit model has the highest loss. Tree-based models perform much better than logit.

Picking the best specification

For each method's best specification we build the lift curve

```
#
```

```

nmethod = length(phatL)
phatBest = matrix(0.0,nrow(testDf),nmethod)
colnames(phatBest) = names(phatL)
for(i in 1:nmethod) {
  nrun = ncol(phatL[[i]])
  lvec = rep(0,nrun)
  print(nrun)
  for(j in 1:nrun) lvec[j] = lossf(testDf$y,phatL[[i]][,j])
  print(lvec)
  imin = which.min(lvec)
  cat("imin: ",imin,"\n")
  phatBest[i] = phatL[[i]][,imin]
  phatBest[i] = phatL[[i]][,1]
}


```

```

dfrac = (1:nrow(testDf))/nrow(testDf)
plot(c(0,1),c(0,1),xlab='% tried',ylab='% of successes',cex=2,type="n", main="Lift curves")
for(i in 1:ncol(phatBest)) {
  temp = liftf(testDf$y,phatBest[,i],dopl=FALSE)
  lines(dfrac,temp,type="l",col=i)
}
abline(0,1,lty=2)
legend("bottomright",legend=names(phatL),
      col=1:nmethod,lty=rep(1,nmethod), cex=0.8,)

```

Lift curves



So, the best methods and specifications are:

- a) Random Forest specification 4 (mtry=18.27567,ntree=1000)

b) Boosting specification 8 (depth=4,n.trees=6000,shrinkage=.01)

Now, we fit models and plot ROC curves for the best rf, boosting and the unique logit model

```
# rf

rf_best = randomForest(y~.,data=trainDf,mtry=18.27567,ntree=1000)

rf_pred_best = predict(rf_best,newdata=testDf,type="prob")[,2]

# gbm

gbm_best = gbm(y~.,data=trainDfB,distribution="bernoulli",
               interaction.depth=4,n.trees=6000,shrinkage=.01)

gbm_pred_best = predict(gbm_best,newdata=testDfB, n.trees=6000,type="response")

#logit

lgfit = glm(y~.,trainDf,family=binomial)


phat = predict(lgfit,testDf,type="response")

## Warning in predict.lm(object, newdata, se.fit, scale = 1, type =
## ifelse(type == : prediction from a rank-deficient fit may be misleading

pred1 <- prediction(rf_pred_best, testDf$y)
pred9 <- prediction(phat, testDf$y)
pred4 <- prediction(gbm_pred_best, testDf$y)

perf1 <- performance(pred1,"tpr","fpr")
perf9 <- performance(pred9,"tpr","fpr")
perf4 <- performance(pred4,"tpr","fpr")

plot(perf9, col=4, main = "ROC curves")
lines(perf4@x.values[[1]], perf4@y.values[[1]], col = 2)
lines(perf1@x.values[[1]], perf1@y.values[[1]], col = 1)
abline(0,1,lty=2)
legend("bottomright",
      legend=c("Logit", "GBM", "RF"),
      col=c("blue","red","black"), lty=c(1,1), cex=0.8)
```



Comparison of variable importance form different models

GLM

```
head(varImp(lgfit), n=20)
```

```
##           Overall
## POP2010    19655854
## POP10_SQMI 26336341
## POP2012    44093012
## POP12_SQMI 26234559
## WHITE      25475496
## BLACK      35977124
## AMERI_ES   61133435
## ASIAN       2861453
## HAWN_PI    58714781
## HISPANIC   40384000
## OTHER      26465806
## MALES      136529434
## AGE_UNDER5 23783434
## AGE_5_9    14516149
## AGE_10_14  10079067
## AGE_15_19  13107504
## AGE_20_24  43589069
## AGE_25_34  13792236
## AGE_35_44  16718007
## AGE_45_54  33068513
```


GBM

```
head(summary(gbm_best, plotit=FALSE), n=20)
```

```
##           var  rel.inf
## BLACK      BLACK 4.869648
## rca_2361    rca_2361 4.784303
## FHH_CHILD  FHH_CHILD 3.867464
## AMERI_ES    AMERI_ES 3.817313
## rca_4529    rca_4529 3.775912
## MED_AGE_M  MED_AGE_M 2.817762
## VACANT      VACANT 2.722560
## rca_2381    rca_2381 2.617724
## rca_5222    rca_5222 2.505626
## MARHH_CHD  MARHH_CHD 2.465957
## rca_1133    rca_1133 2.340757
## rca_4461    rca_4461 1.963703
## rca_2383    rca_2383 1.898341
## rca_6241    rca_6241 1.819618
## rca_4451    rca_4451 1.757335
## ASIAN      ASIAN 1.645390
## PNTCNT_S    PNTCNT_S 1.551606
## rca_6214    rca_6214 1.424837
## AGE_85_UP  AGE_85_UP 1.301156
## OTHER      OTHER 1.298451
```

RF

```
head(importance(rf_best), n=20)
```

```
##           MeanDecreaseGini
## POP2010           8.533410
## POP10_SQMI        11.934628
## POP2012           8.528282
## POP12_SQMI        11.800887
## WHITE             12.934156
## BLACK             26.748412
## AMERI_ES          13.667884
## ASIAN             10.844283
## HAWN_PI           8.149513
## HISPANIC          11.313124
## OTHER             11.740289
## MULT_RACE         10.568144
## MALES             8.752840
## FEMALES           8.355474
## AGE_UNDER5        9.058235
## AGE_5_9           8.917370
## AGE_10_14         9.148620
## AGE_15_19         8.814471
## AGE_20_24        12.397916
## AGE_25_34         9.556742
```

We can observe that the variable importance varies across models. Interestingly, Boosting model, which is the model that has the best predictive capacity, has the largest number of predictors from the competitiveness index matrix.

```
stopCluster(c1)
```